

SPRAWOZDANIE

Zajęcia: Eksploracja i wizualizacja danych
Prowadzący: prof. dr hab. Vasyl Martsenyuk

Laboratorium: 3

Data: 23.03.2023

Temat: "Użycie biblioteki PySpark dla dużych zbiorów danych"

Wariant: 7

Michał Stajerski
Informatyka II stopień,
stacjonarne, semestr 3,
Gr. 1

<https://github.com/MichalStajerski/eiwd>

1. Polecenie

Celem zajęć jest eksploracja i wizualizacja danych z użyciem API Spark przez bibliotekę pyspark. Kolejnym etapem zajęć jest pobieranie danych z pliku CSV lokalnie oraz z

użyciem URL. Podstawowe kroki eksploracji przedstawione są w pliku Zajecie3_Spark.ipynb. Dane są określone wariantem z zadania 1, które zostały pobrane ze strony https://ghdx.healthdata.org/ihme_data. Wariant wybrany w zadaniu jest wariant 7:

Global Burden of Disease Study 2019 (GBD 2019) Smoking Tobacco Use Prevalence 1990-2019

2. Zadania

1 - Użycie PySpark w celu eksploracji Big Data

1.1 - Konfigurowanie środowiska w Anaconda lub Google Colab

Do wykonywania zadań został wybrany Google Colaboratory

```
[88] # Polecenie do zainstalowania pakietów pyspark i py4j:
```

```
! pip install pyspark==3.0.1 py4j==0.10.9
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: pyspark==3.0.1 in /usr/local/lib/python3.10/dist-packages (3.0.1)
Requirement already satisfied: py4j==0.10.9 in /usr/local/lib/python3.10/dist-packages (0.10.9)
```

1.2 - Sesja Spark

1.3 - Tworzenie SparkSession

```
[89] from pyspark.sql import SparkSession
      spark = SparkSession.builder\
            .master("local[*]")\
            .appName('PySpark_Tutorial')\
            .getOrCreate()
```

1.4 - Czytanie danych

```
[90] csv_file = '/content/IHME_GBD_2019_SMOKING_TOB_1990_2019_NUM_SMOKERS_Y2021M05D27.CSV'
      df = spark.read.csv(csv_file)
```

1.5 - Pobieranie danych za pomocą URL

```
[91] from pyspark import SparkFiles

      spark.sparkContext.addFile('https://storage.covid19datahub.io/level/1.csv')

      df = spark.read.csv(SparkFiles.get("1.csv"), header=True)
```

2 - Strukturyzacja danych za pomocą schematu Spark Kod do odczytu danych w formacie pliku CSV:

```
[92] data = spark.read.csv(  
    '/content/IHME_GBD_2019_SMOKING_TOB_1990_2019_NUM_SMOKERS_Y2021M05D27.CSV',  
    sep=',',  
    header=True,  
)  
data.printSchema()
```

```
root  
|-- measure_name: string (nullable = true)  
|-- location_id: string (nullable = true)  
|-- location_name: string (nullable = true)  
|-- sex_id: string (nullable = true)  
|-- sex_name: string (nullable = true)  
|-- age_group_id: string (nullable = true)  
|-- age_group_name: string (nullable = true)  
|-- year_id: string (nullable = true)  
|-- val: string (nullable = true)  
|-- upper: string (nullable = true)  
|-- lower: string (nullable = true)
```

Precyzowanie struktury danych

```
[93] from pyspark.sql.types import *
data_schema = [
    StructField('measure_name', StringType(), True),
    StructField('location_id', IntegerType(), True),
    StructField('location_name', StringType(), True),
    StructField('sex_id', IntegerType(), True),
    StructField('sex_name', StringType(), True),
    StructField('age_group_id', IntegerType(), True),
    StructField('age_group_name', StringType(), True),
    StructField('year_id', IntegerType(), True),
    StructField('val', DoubleType(), True),
    StructField('upper', DoubleType(), True),
    StructField('lower', DoubleType(), True),
]

final_struct = StructType(fields = data_schema)
data = spark.read.csv(
    '/content/IHME_GBD_2019_SMOKING_TOB_1990_2019_NUM_SMOKERS_Y2021M05D27.CSV',
    sep=',',
    header=True,
    schema=final_struct
)
data.printSchema()
```

```
root
|-- measure_name: string (nullable = true)
|-- location_id: integer (nullable = true)
|-- location_name: string (nullable = true)
|-- sex_id: integer (nullable = true)
|-- sex_name: string (nullable = true)
|-- age_group_id: integer (nullable = true)
|-- age_group_name: string (nullable = true)
|-- year_id: integer (nullable = true)
|-- val: double (nullable = true)
|-- upper: double (nullable = true)
|-- lower: double (nullable = true)
```

3 - Różne metody kontroli danych

3.1 - schema(): Ta metoda zwraca schemat danych (ramka danych).

```
data.schema
StructType(List(StructField(measure_name,StringType,true),StructField(location_id,IntegerType,true),StructField(location_name,StringType,true),St
```

3.2 - dtypes zwraca listę krotek z nazwami kolumn i typami danych.

→

3.3 - head(n) zwraca n wierszy jako listę.

```
[96] data.head(3)
```

```
[Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=1, sex_name='Male', age_group_id=29, age_group_name='15+ years', year_id=1990, val=803101467.1, upper=809622101.0, lower=795908635.8),
Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=2, sex_name='Female', age_group_id=29, age_group_name='15+ years', year_id=1990, val=189148834.0, upper=193092888.7, lower=185559469.9),
Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=3, sex_name='Both', age_group_id=29, age_group_name='15+ years', year_id=1990, val=992250301.2, upper=1000161258.0, lower=984788043.8)]
```

3.4 - `show()` domyślnie wyświetla pierwsze 20 wierszy, a także przyjmuje liczbę jako parametr określający ich liczbę

data.show()

only showing top 20 rows

3.5 first() zwraca pierwszy wiersz danych.

```
[98] data.first()
```

```
Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=1, sex_name='Male', age_group_id=29, age_group_name='15+ years', year_id=1990, val=803101467.1, upper=809622101.0, lower=795908635.8)
```

3.6 - take(n) zwraca pierwsze n wierszy

```
[99] data.take(5)
```

```
[Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=1, sex_name='Male', age_group_id=29, age_group_name='15+ years', year_id=1990, val=803101467.1, upper=809622101.0, lower=795908635.8),
 Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=2, sex_name='Female', age_group_id=29, age_group_name='15+ years', year_id=1990, val=189148834.0, upper=193092888.7, lower=185559469.9),
 Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=3, sex_name='Both', age_group_id=29, age_group_name='15+ years', year_id=1990, val=992250301.2, upper=1000161258.0, lower=984788043.8),
 Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=1, sex_name='Male', age_group_id=29, age_group_name='15+ years', year_id=1991, val=813897216.4, upper=820033926.0, lower=806951447.9),
 Row(measure_name='Number of Smokers', location_id=1, location_name='Global', sex_id=2, sex_name='Female', age_group_id=29, age_group_name='15+ years', year_id=1991, val=190537545.1, upper=194424928.7, lower=186974424.5)]
```

3.7 - describe() oblicza niektóre wartości statystyczne dla kolumn liczbowych.

```
[100] data.describe()
```

```
DataFrame[summary: string, measure_name: string, location_id: string, location_name: string, sex_id: string, sex_name: string, age_group_id: string, age_group_name: string, year_id: string, val: string, upper: string, lower: string]
```

3.8 - columns zwraca listę zawierającą nazwy kolumn.

```
[101] data.columns
```

```
['measure_name',
 'location_id',
 'location_name',
 'sex_id',
 'sex_name',
 'age_group_id',
 'age_group_name',
 'year_id',
 'val',
 'upper',
 'lower']
```

3.9 - count() zwraca całkowitą liczbę wierszy w zestawie danych

```
[▶] data.count()
```

```
↗ 20970
```

3.10 - differ() to liczba odmiennych wierszy w używanym zbiorze danych.

```
[103] data.differ()
```

3.11 - printSchema() wyświetla schemat danych.

```
[104] data.printSchema()
```

```
root
|-- measure_name: string (nullable = true)
|-- location_id: integer (nullable = true)
|-- location_name: string (nullable = true)
|-- sex_id: integer (nullable = true)
|-- sex_name: string (nullable = true)
|-- age_group_id: integer (nullable = true)
|-- age_group_name: string (nullable = true)
|-- year_id: integer (nullable = true)
|-- val: double (nullable = true)
|-- upper: double (nullable = true)
|-- lower: double (nullable = true)
```

4- Manipulacja kolumnami

4.1 - Dodawanie kolumny:

```
data = data.withColumn('location_name_copy', data.location_name)
data.show(5)
```

| location_name | location_id | location_name | sex_id | sex_name | age_group_id | age_group_name | year_id | val | upper | lower | location_name_copy |
|---------------|-------------|---------------|--------|----------|--------------|----------------|---------|---------------|---------------|---------------|--------------------|
| tokers | 1 | Global | 1 | Male | 29 | 15+ years | 1990 | 8.031014671E8 | 8.09622101E8 | 7.959086358E8 | Global |
| tokers | 1 | Global | 2 | Female | 29 | 15+ years | 1990 | 1.89148834E8 | 1.930928887E8 | 1.855594699E8 | Global |
| tokers | 1 | Global | 3 | Both | 29 | 15+ years | 1990 | 9.922503012E8 | 1.000161258E9 | 9.847880438E8 | Global |
| tokers | 1 | Global | 1 | Male | 29 | 15+ years | 1991 | 8.138972164E8 | 8.20033926E8 | 8.069514479E8 | Global |
| tokers | 1 | Global | 2 | Female | 29 | 15+ years | 1991 | 1.905375451E8 | 1.944249287E8 | 1.869744245E8 | Global |

top 5 rows

4.2 - Aktualizacja kolumny:

```
data = data.withColumnRenamed('location_name_copy', 'location_name_changed')
data.show(5)
```

| location_id | location_name | sex_id | sex_name | age_group_id | age_group_name | year_id | val | upper | lower | location_name_changed |
|-------------|---------------|--------|----------|--------------|----------------|---------|---------------|---------------|---------------|-----------------------|
| 1 | Global | 1 | Male | 29 | 15+ years | 1990 | 8.031014671E8 | 8.09622101E8 | 7.959086358E8 | Global |
| 1 | Global | 2 | Female | 29 | 15+ years | 1990 | 1.89148834E8 | 1.930928887E8 | 1.855594699E8 | Global |
| 1 | Global | 3 | Both | 29 | 15+ years | 1990 | 9.922503012E8 | 1.000161258E9 | 9.847880438E8 | Global |
| 1 | Global | 1 | Male | 29 | 15+ years | 1991 | 8.138972164E8 | 8.20033926E8 | 8.069514479E8 | Global |
| 1 | Global | 2 | Female | 29 | 15+ years | 1991 | 1.905375451E8 | 1.944249287E8 | 1.869744245E8 | Global |

5 rows

4.3 Upuszczanie kolumny:

-

```
[107] data = data.drop('location_name_changed')
data.show(5)
```

| measure_name | location_id | location_name | sex_id | sex_name | age_group_id | age_group_name | year_id | val | upper | lower |
|-------------------|-------------|---------------|--------|----------|--------------|----------------|---------|---------------|---------------|---------------|
| Number of Smokers | 1 | Global | 1 | Male | 29 | 15+ years | 1990 | 8.031014671E8 | 8.09622101E8 | 7.959086358E8 |
| Number of Smokers | 1 | Global | 2 | Female | 29 | 15+ years | 1990 | 1.89148834E8 | 1.930928887E8 | 1.855594699E8 |
| Number of Smokers | 1 | Global | 3 | Both | 29 | 15+ years | 1990 | 9.922503012E8 | 1.000161258E9 | 9.847880438E8 |
| Number of Smokers | 1 | Global | 1 | Male | 29 | 15+ years | 1991 | 8.138972164E8 | 8.20033926E8 | 8.069514479E8 |
| Number of Smokers | 1 | Global | 2 | Female | 29 | 15+ years | 1991 | 1.905375451E8 | 1.944249287E8 | 1.869744245E8 |

only showing top 5 rows

5- Radzenie sobie z brakującymi wartościami

```
from pyspark.sql import functions as f
# Usuń wiersze z brakującymi wartościami w dowolnej z kolumn
data.na.drop()
# Zastąp brakujące wartości za pomocą średniej
data.na.fill(data.select(f.mean(data['age_group_id'])).collect()[0][0])
# Zastąp brakujące wartości nowymi
data.na.replace(old_value, new_value)
```

DataFrame[measure_name: string, location_id: int, location_name: string, sex_id: int, sex_name: string, age_group_id: int, age_group_name: string, year_id: int, val: double, upper: double, lower: double]

6- Pobieranie danych

6.1 - Select

```
[109] # wybór jednej kolumny
data.select('year_id').show(5)
```

| year_id |
|---------|
| 1990 |
| 1990 |
| 1990 |
| 1991 |
| 1991 |

only showing top 5 rows


```
# wybór kilku kolumn
data.select(['location_name', 'sex_name', 'year_id']).show(5)
```

```
+-----+-----+-----+
|location_name|sex_name|year_id|
+-----+-----+-----+
|      Global|    Male|   1990|
|      Global|   Female|   1990|
|      Global|    Both|   1990|
|      Global|    Male|   1991|
|      Global|   Female|   1991|
+-----+-----+-----+
```

only showing top 5 rows

6.2 - Filter

```
[111] from pyspark.sql.functions import col
data.filter( (col('year_id') >= 2000) & (col('sex_name') == 'Female') ).show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|measure_name|location_id|location_name|sex_id|sex_name|age_group_id|age_group_name|year_id|val|upper|lower|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Number of Smokers|1|Global|2|Female|29|15+ years|2000|2.03389244E8|2.070119921E8|2.000443639E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|2001|2.043228487E8|2.078410984E8|2.011225869E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|2002|2.051251323E8|2.086136115E8|2.019432714E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|2003|2.05852521E8|2.094216076E8|2.026372326E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|2004|2.064882331E8|2.099930308E8|2.031999303E8|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

only showing top 5 rows

6.3 - Between

```
data.filter(data.year_id.between(1995, 2000)).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|measure_name|location_id|location_name|sex_id|sex_name|age_group_id|age_group_name|year_id|val|upper|lower|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Number of Smokers|1|Global|1|Male|29|15+ years|1995|8.433043019E8|8.490080491E8|8.3750097E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|1995|1.963544335E8|2.002139588E8|1.930128599E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|1995|1.039658735E9|1.047062623E9|1.032850499E9|
|Number of Smokers|1|Global|1|Male|29|15+ years|1996|8.478849471E8|8.536353541E8|8.42071989E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|1996|1.980633863E8|2.018466552E8|1.946320182E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|1996|1.045948333E9|1.053276972E9|1.039024799E9|
|Number of Smokers|1|Global|1|Male|29|15+ years|1997|8.516471918E8|8.573761175E8|8.457102366E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|1997|1.996329437E8|2.032874241E8|1.962767139E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|1997|1.051280135E9|1.05835932E9|1.044430801E9|
|Number of Smokers|1|Global|1|Male|29|15+ years|1998|8.550585577E8|8.608097687E8|8.491967263E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|1998|2.010668729E8|2.048237827E8|1.978497379E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|1998|1.056125431E9|1.062905206E9|1.049476558E9|
|Number of Smokers|1|Global|1|Male|29|15+ years|1999|8.582672178E8|8.639698017E8|8.525235784E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|1999|2.023213713E8|2.059460479E8|1.989813106E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|1999|1.060588589E9|1.067665003E9|1.053944039E9|
|Number of Smokers|1|Global|1|Male|29|15+ years|2000|8.6154577E8|8.67412168E8|8.560329846E8|
|Number of Smokers|1|Global|2|Female|29|15+ years|2000|2.03389244E8|2.070119921E8|2.000443639E8|
|Number of Smokers|1|Global|3|Both|29|15+ years|2000|1.064935014E9|1.071795767E9|1.058216329E9|
|Number of Smokers|4|Southeast Asia, E...|1|Male|29|15+ years|1995|3.779077105E8|3.82356775E8|3.731279345E8|
|Number of Smokers|4|Southeast Asia, E...|2|Female|29|15+ years|1995|2.964856863E7|3.221889337E7|2.751987308E7|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

only showing top 20 rows

6.4 When

```
[113] data.select('location_name', 'sex_name',
  f.when(data.year_id == '2000', 1).otherwise(0)
).show(5)
```

```

+-----+-----+-----+
|location_name|sex_name|CASE WHEN (year_id = 2000) THEN 1 ELSE 0 END|
+-----+-----+-----+
|      Global|   Male|                                           0|
|      Global| Female|                                           0|
|      Global|   Both|                                           0|
|      Global|   Male|                                           0|
|      Global| Female|                                           0|
+-----+-----+-----+
only showing top 5 rows

```

6.5 - Like

```
[114] data.select('location_name', data.location_name.rlike('^[G,P]').alias('location_name zaczyna sie na litere G lub P')).distinct().show()
```

```

+-----+-----+
|location_name|location_name zaczyna sie na litere G lub P|
+-----+-----+
|      Cuba|false|
|Mauritania|false|
|Djibouti|false|
|Slovenia|false|
|Sub-Saharan Africa|false|
|      Malawi|false|
|United Kingdom|false|
|      Pakistan|true|
|      Botswana|false|
|      Madagascar|false|
|      Australia|false|
|United States of ...|false|
|      Ghana|true|
|      Tokelau|false|
|      Belarus|false|
|Bolivia (Plurinat...|false|
|Dominican Republic|false|
|      Paraguay|true|
|      Croatia|false|
|      Ukraine|false|
+-----+-----+
only showing top 20 rows

```

6.6 - GroupBy

```
[115] data.select(['val', 'upper', 'lower']).groupBy('val').mean().show()
```

| val | avg(val) | avg(upper) | avg(lower) |
|---------------|---------------|---------------|--------------|
| 134007.9694 | 134007.9694 | 139677.191 | 128260.0995 |
| 8270.464532 | 8270.464532 | 9504.142112 | 7096.286328 |
| 11829.30213 | 11829.30213 | 13621.19377 | 10115.58969 |
| 35698.65734 | 35698.65734 | 37600.2659 | 33870.85627 |
| 46355.98657 | 46355.98657 | 49432.5945 | 43154.77478 |
| 13869.07463 | 13869.07463 | 15706.97617 | 12179.06027 |
| 20742.10263 | 20742.10263 | 22062.18255 | 19357.08588 |
| 7423.071697 | 7423.071697 | 8228.304212 | 6694.14611 |
| 7448.841156 | 7448.841156 | 8213.962619 | 6706.265256 |
| 9104.794872 | 9104.794872 | 10096.50032 | 8162.279112 |
| 2887.470373 | 2887.470373 | 3818.180825 | 2189.147345 |
| 4744.936929 | 4744.936929 | 5932.17345 | 3729.904017 |
| 458.0840221 | 458.0840221 | 606.3972438 | 344.4231654 |
| 3185.522508 | 3185.522508 | 3507.963054 | 2861.54479 |
| 3986.842432 | 3986.842432 | 4810.125966 | 3191.818067 |
| 7677.332734 | 7677.332734 | 8179.958609 | 7109.126253 |
| 12568.07874 | 12568.07874 | 13358.02037 | 11766.13512 |
| 4478.782127 | 4478.782127 | 4815.00638 | 4148.004533 |
| 441.7884447 | 441.7884447 | 565.9563718 | 331.323542 |
| 8.313872544E8 | 8.313872544E8 | 8.372931128E8 | 8.24949648E8 |

only showing top 20 rows

6.7 - Agregacja

```
[116] from pyspark.sql import functions as f
```

```
data.filter((col('val') >= 1000000) & (col('val') <= 5000000))\
.groupBy("location_name") \
.agg(f.min("val").alias("from"),
f.max("val").alias("to"),
f.min("val").alias("minimum vaccinated"),
f.max("val").alias("maximum vaccinated"),
f.avg("val").alias("average vaccinated"),
f.min("val").alias("minimum_economic_support_index"),
f.max("val").alias("maximum_economic_support_index"),
f.avg("val").alias("average_economic_support_index"),
).show(truncate=False)
```


| location_name | from | to | minimum vaccinated | maximum vaccinated | average vaccinated | minimum_economic_support_index | maximum_econor |
|-------------------|-------------|-------------|--------------------|--------------------|--------------------|--------------------------------|----------------|
| Côte d'Ivoire | 1023886.493 | 2165483.532 | 1023886.493 | 2165483.532 | 1601788.6232857148 | 1023886.493 | 2165483.532 |
| Yemen | 1030535.291 | 3657898.202 | 1030535.291 | 3657898.202 | 1993066.6792333333 | 1030535.291 | 3657898.202 |
| Sweden | 1067123.166 | 1644302.372 | 1067123.166 | 1644302.372 | 1308458.0137666664 | 1067123.166 | 1644302.372 |
| Republic of Korea | 1063356.593 | 1179678.716 | 1063356.593 | 1179678.716 | 1147212.8582666668 | 1063356.593 | 1179678.716 |
| Philippines | 2214973.74 | 3144082.409 | 2214973.74 | 3144082.409 | 2863151.9060000004 | 2214973.74 | 3144082.409 |
| Malaysia | 2942804.951 | 4978849.254 | 2942804.951 | 4978849.254 | 4206188.935574467 | 2942804.951 | 4978849.254 |
| Turkey | 3198400.064 | 4940216.419 | 3198400.064 | 4940216.419 | 4282831.987909092 | 3198400.064 | 4940216.419 |
| Malawi | 1019364.093 | 1164664.74 | 1019364.093 | 1164664.74 | 1083080.244666667 | 1019364.093 | 1164664.74 |
| Iraq | 2135606.051 | 4941969.609 | 2135606.051 | 4941969.609 | 3431766.7088846145 | 2135606.051 | 4941969.609 |
| Cambodia | 1089526.703 | 2373559.505 | 1089526.703 | 2373559.505 | 1725967.798783333 | 1089526.703 | 2373559.505 |
| Afghanistan | 1013802.294 | 2257180.971 | 1013802.294 | 2257180.971 | 1605877.0142962963 | 1013802.294 | 2257180.971 |
| Jordan | 1015421.261 | 2817093.164 | 1015421.261 | 2817093.164 | 1765783.7067692312 | 1015421.261 | 2817093.164 |
| Sudan | 1210513.281 | 2689229.553 | 1210513.281 | 2689229.553 | 1917168.4488166673 | 1210513.281 | 2689229.553 |
| Greece | 1269978.149 | 3921930.694 | 1269978.149 | 3921930.694 | 2422811.3991111116 | 1269978.149 | 3921930.694 |
| Sri Lanka | 2073078.052 | 2606806.575 | 2073078.052 | 2606806.575 | 2329932.261716666 | 2073078.052 | 2606806.575 |
| Algeria | 2473253.484 | 4979791.166 | 2473253.484 | 4979791.166 | 3636862.253909091 | 2473253.484 | 4979791.166 |
| Slovakia | 1156863.894 | 1331143.078 | 1156863.894 | 1331143.078 | 1260198.3315666674 | 1156863.894 | 1331143.078 |
| Argentina | 2982430.438 | 4846195.904 | 2982430.438 | 4846195.904 | 3803538.7871000012 | 2982430.438 | 4846195.904 |
| Belgium | 1007284.838 | 2783487.184 | 1007284.838 | 2783487.184 | 1656247.782831326 | 1007284.838 | 2783487.184 |
| Angola | 1002162.849 | 1519341.656 | 1002162.849 | 1519341.656 | 1212375.1044736842 | 1002162.849 | 1519341.656 |

only showing top 20 rows

7- wizualizacja danych

```
[117] from matplotlib import pyplot as plt

currency_df = data.select(['location_name',
                           'val',
                           'upper',
                           'lower'])
currency_df.groupby('location_name')\
    .mean()\
    .toPandas()
ind = list(range(12))
ind.pop(6)
currency_df.iloc[ind,:].plot(kind='bar', x='location_name', y=currency_df.columns.tolist()[1:],
figsize=(12, 6), ylabel='People', xlabel='iso_currency')
plt.show()
```

8- Zapisywanie/zapisywanie danych do pliku

```
[80] # CSV
data.write.csv('dataset.csv')
# JSON
data.write.save('dataset.json', format='json')
# Parquet
data.write.save('dataset.parquet', format='parquet')
```

```
[87] # Zapisywanie wybranych kolumn
# CSV
data.select(['location_name', 'sex_name', 'year_id', 'val'])\
.write.csv('dataset_columns.csv')
# JSON
data.select(['location_name', 'sex_name', 'year_id', 'val'])\
.write.save('dataset_columns.json', format='json')
# Parquet
data.select(['location_name', 'sex_name', 'year_id', 'val'])\
.write.save('dataset_columns.parquet', format='parquet')
```


3. Wnioski

Na podstawie otrzymanego wyniku można stwierdzić, że biblioteka PySpark używana jest do przetwarzania dużych zbiorów danych. Używana podczas zajęć biblioteka pozwala na wczytywanie danych z różnych formatów: CSV, JSON, Parquet, jak również za pomocą URL. Biblioteka PySpark pozwala manipulować kolumnami: dodawać kolumnę, zmienić nazwę kolumnie, oraz usunąć kolumnę. PySpark i PySpark SQL zapewniają szeroki zakres metod i funkcji do łatwego wyszukiwania danych.